

===== Relational Algebra =====

THETA JOIN:

$$R_1 \bowtie_C R_2 = \sigma_C(R_1 \times R_2)$$

INTERSECT:

$$R \cap S = R - (R - S)$$

UNION, DIFFERENCE, INTERSECT Operators:

- Schemas must be the same
- No duplicates remain

DIVISION:

$$R/S = \pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$$

R/S is the largest relation T such that $T \times S \subseteq R$

===== SQL =====

General statement:

SELECT attributes, aggregates

FROM relations(tables)

WHERE conditions

GROUP BY attributes

HAVING conditions on aggregates

ORDER BY attributes, aggregates

Subqueries

- in WHERE clause
 - Considered as a regular relation
 - One-attribute one-tuple relation -> use like a 'value'
- In FROM clause
 - Considered as a regular relation
 - Must be renamed to a new table name

Evaluation order:

FROM → WHERE → GROUP BY → HAVING → ORDER BY → SELECT

Aggregates

- Sum, Count, Avg, Min, Max, ...

Set Operators:

INTERSECT, UNION, EXCEPT

- Follow set semantics and remove duplicates
- To keep duplicates: UNION ALL, INTERSECT ALL, EXCEPT ALL

Set membership

- IN, NOT IN

Set comparison operator

- ALL, < SOME, = SOME, ... etc.

INSERT, DELETE, UPDATE

- Insertion: INSERT INTO <Relation> <Tuples>
- Deletion: DELETE FROM <R> WHERE <Condition>
- Update: Update R
 SET $A_1 = V_1, A_2 = V_2, \dots, A_n = V_n$
 WHERE <Condition>

===== Database Integrity =====

Key Constraints

CREATE TABLE <name> (

...,

PRIMARY KEY(dept, cnum, sec),

UNIQUE(dept, cnum, instructor))

- Change E.A value to NULL or default value
- CASCADE
 - On deletion of S: delete referencing tuples in E
 - On update of S.A: change E.A to the new S.A

Referential Integrity (Foreign Key)

- E.A references S.A
 - E.A: referencing attribute / foreign key
 - S.A: referenced attribute
- CREATE TABLE <name> (
 - ...,
 - sid INTEGER REFERENCES Student(sid),
 - FOREIGN KEY (dept, cnum, sec) REFERENCES Class(dept, cnum, sec)
- Referenced attributes must be PRIMARY KEY or UNIQUE
- RI Violation
 - Default: not allowed
 - System rejects the statement
 - Always insert/update S first
 - ON DELETE/UPDATE SET NULL/SET DEFAULT/CASCADE
 - Added on Referencing attributes declaration
 - SET NULL/SET DEFAULT

Check Constraints

- CRATE TABLE Enroll (
 - dept CHAR(2), cnum INT, unit INT,
 - title VARCHAR(50),
 - CHECK (cnum < 600 AND unit < 10))

Triggers

CREATE Trigger <name>

<event>

<referencing clause>

WHEN (<condition>)

<action>

<event>

- BEFORE | AFTER INSERT/DELETE/UPDATE [OF A1, A2, ..., An] ON R

<referencing clause>

- REFERENCING OLD | NEW TABLE | ROW AS <var>, ...
- FOR EACH ROW | STATEMENT

<action>

- Any SQL statement

===== Views and Authorization =====

Views

CREATE VIEW <name> AS

<Query>

Authorization

- GRANT <privileges> ON <R> TO <user> [WITH GRANT OPTION]
- REVOKE <privileges> ON <R> FROM <user> [CASCADE | RESTRICT]

===== Misc =====

Logical Implies in SQL

$$p \rightarrow q \equiv \neg p \vee q$$

===== Files and Disks =====

Access time = (seek time) + (rotational delay) + (transfer time)

- Seek time: time to find the target track
 - Typical average seek time: 10 ms
- Rotational delay: time to rotate to the target sector
 - For 6000 RPM, average rotational delay = $0.5 * (1\text{min}/6000) = 0.5 * 60\text{sec}/6000 = 5\text{ ms}$
- Transfer Time
 - Time to read one block
 - For example, 6000 RPM, 1000 sector/track, 1KB/sector
 - Read a track, rotate a circle: $1\text{min}/6000 = 10\text{ ms/track}$
 - Read one sector(block): $(10\text{ms/track}) / (1000\text{sector/track}) = \mathbf{0.01\text{ms/sector}}$
 - Transfer rate: $1\text{KB}/(0.01\text{ms/sector}) = \mathbf{100\text{MB/s}}$

===== B+ Trees =====

Insertion:

- Leaf node overflow
 - The first key of the new node is copied to the parent
- Non-leaf node overflow
 - The middle key is moved to the new parent

	MaxPtrs	MaxKeys	MinPtrs	MinKeys
Non-leaf Non-root	n	n-1	$\lceil n/2 \rceil$	$\lceil n/2 \rceil - 1$
Leaf Non-root	n	n-1	$\lceil (n+1)/2 \rceil$	$\lceil (n-1)/2 \rceil$
Root	n	n-1	2	1